

Caller-saved, callee-saved registers

- when you are writing a function in assembly, that function will be called by another function, so your function is the "callee".

_f:

==== } able to overwrite the
"caller-saved registers"

call g } here, I am the caller,
so I am responsible for
saving & restoring whatever
caller-saved registers I need.
} back in the "callee" part.

How do I save a caller-saved register?

- ① Move its contents into a callee-saved register
- ② Push its contents on the stack.



mov \$24, %rcx
push %rcx #^{caller saved} on stack
call -q
pop %rcx #^{restore from} #^{stack}
add \$12, %rcx

Saving multiple registers?

- pop in reverse order of push.

```
push %rcx  
push %rdx  
call -g  
pop %rdx  
pop %rcx
```

Important: push and pop only work on 64-bit values.

A callee-saved register must be saved before you use it and restored after you're done with it.

- caller expects these registers to be unchanged.

if:

push %rbx # save %rbx

mov \$100, %ebx # callee
saved

done using %ebx

pop %rbx # restoring it.

Typically, the callee-saved registers you are going to use are saved at the beginning of the function (after the "push %rbp" and "mov %rsp, %rbp").

The caller-saved registers that you need to save are typically saved before a function call and restored after the call returns.

Pointers in assembly

C code:

```
int foo(int *p)
```

```
{  
    *p = *p * 2;  
}
```

```
int main()
```

```
{  
    int x = 20;  
    foo(&x);  
    printf("%d\n", x); // 40  
}
```

passed in
%rdi

In assembly:

```
- foo:    push %rbp  
          mov %rsp, %rbp  
          # p is in %rdi  
          imul $2, (%rdi)  
          # *p = *p * 2  
          pop %rbp  
          ret
```

Passing parameters, revisited

Mac/Linux: First 6 parameters
are passed in registers,
rest passed on the
stack in reverse order.

Windows: First 4 in registers,
rest on stack in
reverse order.

See the calling convention
document on Brightspace.

C code:

int bar(int a, int b, int c,
10 params } int d, int e, int f,
so 4
on stack
(for Mac)
and Linux
int g, int h, int i,
int j);

int main()

{ int x;
x=bar(1, 2, 3, 4, 5, 6, 7, 8, 9, 10);

}

How does bar access - the
last 4 parameters on the
stack?

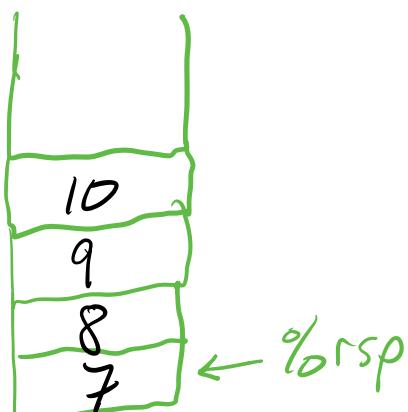
- main() pushes the last 4 parameters in reverse order:

- main :

push \$10
push \$9
push \$8
push \$7

} all 64-bits

What does the stack look like now?

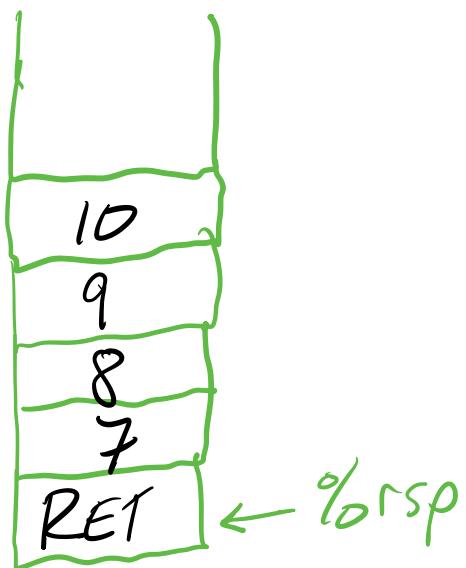


Then main() contains a call instruction:

call -bar

- pushes the return address
onto the stack, jumps
to bar.

So, the stack looks like:



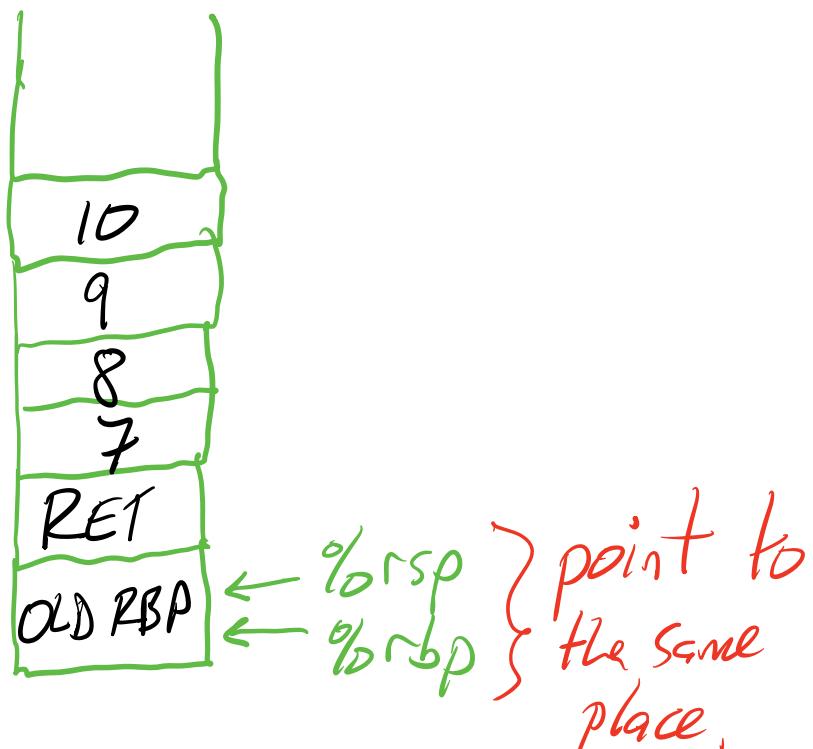
Now the code for bar()
starts executing:

-bar:

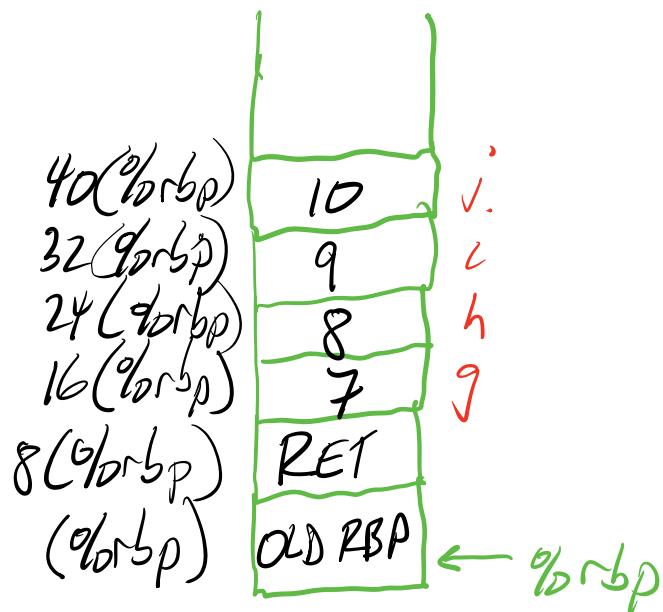
push %rbp
mov %rsp, %rbp

%rbp is callee-saved

Now the stack looks like:



We access parameters on
the stack using offsets
from %rbp

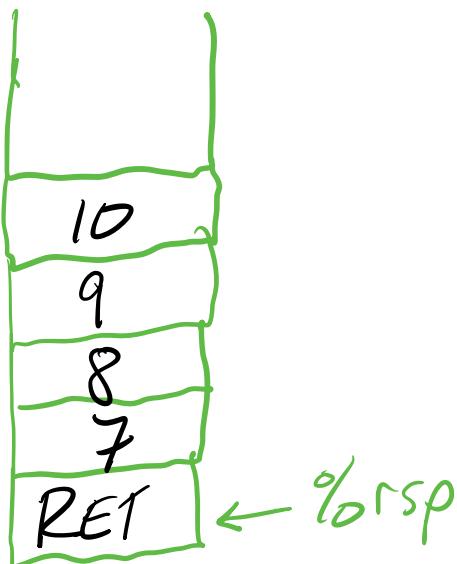


Suppose bar adds those last 4
parameters together.
Each parameter is actually
4 bytes, having been
declared as "int".

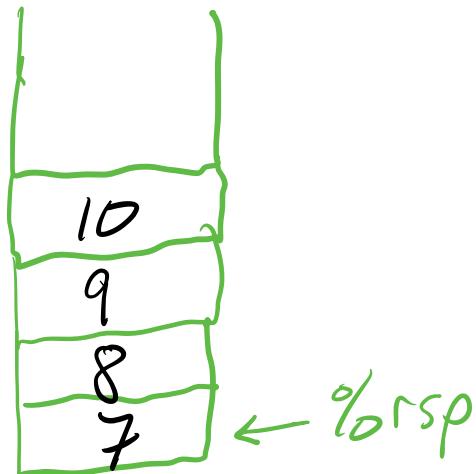
```
movl 16(%rbp), %eax    #eax = q  
addl 24(%rbp), %eax    #eax += h  
addl 32(%rbp), %eax    #eax += i  
addl 40(%rbp), %eax    #eax += j
```

we're done in bar, need to
restore the old %rbp value and
return back to main().

pop %rbp



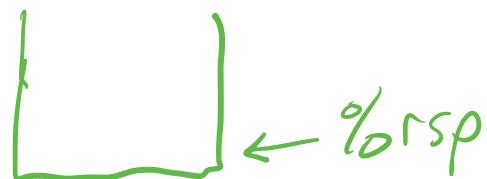
ret # return instruction



Now we're back in main()!

- it needs to remove those 4 parameters from the stack.
- how? It could do 4 pops, but it's more efficient just to add some constant to %rsp.

addq \$32, %rsp # adding 32 to %rsp



The stack is now in the same
state it was before main()
started the call to bar.