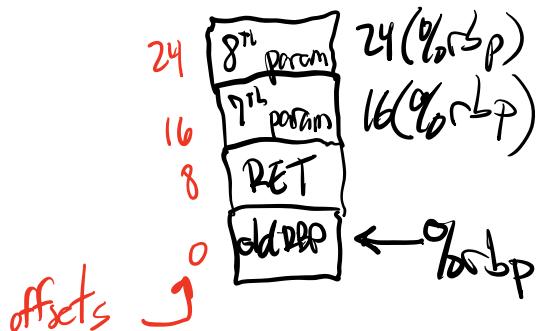


The Stack

MacOS + Linux



Windows

Calling convention says that an extra 32 bytes must be allocated on the stack before a function call.

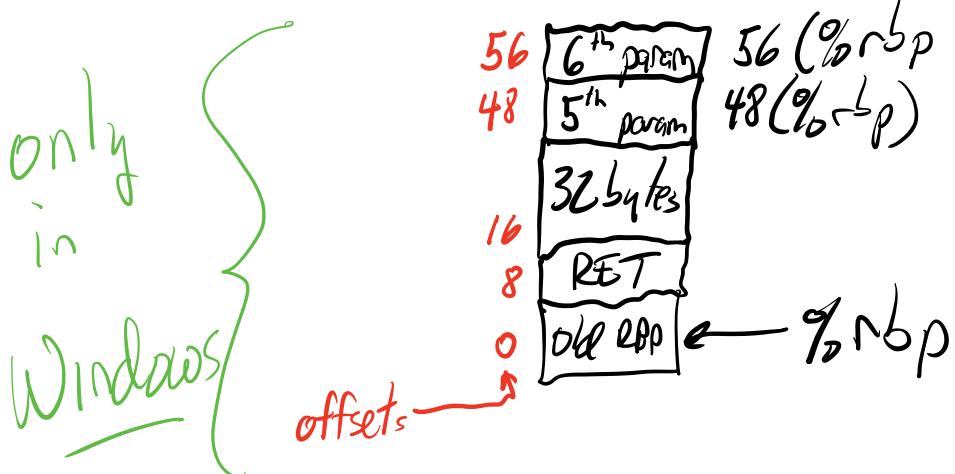
- after any parameters are pushed.

push %rsi } pushing parameters
 push %r8 } (6th & 5th params)
 subq \$32, %rsp ← required by Microsoft

call foo ← Function Call

can be combined as { addq \$32, %rsp ← remove the 32 bytes
 add \$16, %rsp ← remove the parameters
 addq \$48, %rsp

Within the function, the parameter offsets will be different than in macOS:



A register I haven't mentioned yet:

%rip - instruction pointer
(aka "program counter")

The %rip register always points to the next instruction in memory to execute.

- after each mov, add, etc. instruction (not a jump or call), %rip is automatically incremented by the size of the instruction.

- so that it points to the next instruction in memory.

- for a jump, conditional jump, or call instruction, the address of the target instruction is written into %rip.
 - for a conditional jump, this only happens if the condition is satisfied.
- for a return instruction, "ret", the return address is popped off the stack and written to %rip

Allocating local variables on the stack:

C: void f()
{ int x = 10;
int y = 15;
:
}

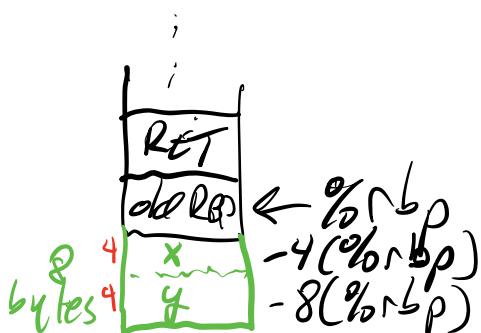
Assembly:

local variables are allocated below where %rbp points:

-f:

push %rbp
mov %rsp, %rbp

allocates space for X and Y.
→
subq \$8, %rsp
movl \$10, -4(%rbp)
movl \$15, -8(%rbp)
;



→
remove the space for the local variables
addq \$8, %rsp
pop %rbp
ret

Important: If you are calling a C library function (`printf`, `malloc`, etc.), keep `%rsp` aligned to 16-byte boundaries

- subtract/add from `%rsp` in multiples of 16.

- this includes when pushing. For example, if you only need to save one register (say, `%rcx`) before calling `printf`, do:

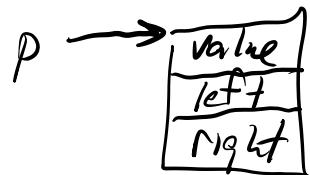
*decrements
`%rsp` by a total
of 16*

```
push %rcx      # pushes 8 bytes
subq $8,%rsp   # adjust %rsp by 8 more.
:
call .printf
addq $8,%rsp   # adjust stack back
pop %rcx       # restore %rcx
```

- This is only if you are calling a C function (anywhere in your program)

Revisiting structs in C

```
typedef struct node {  
    int value;  
    struct node *left;  
    struct node *right;  
} NODE;
```



How do you figure out the offset from the start of the struct for each field?

- it's up to the Compiler,

- ① Write some code, compile it into assembly ("gcc -S ..."), and read the assembly code to figure it out:

```
int main()
{ NODE *p = malloc(sizeof(NODE));
```

```
    int x = p->value;
    NODE *q = p->left;
    q = p->right;
```

{

} could compile
this into assembly
and then need
the assembly code
to see what
the offsets are.

- ② Easier solution: Take addresses of the fields of the NODE p points to and subtract the value of P from them:

```

int main()
{
    NODE *p = malloc(sizeof(NODE));
    printf("offset of p->value = %ld\n",
        (void *) &(p->value) - (void *)p );
    to stop the compiler  
from complaining
    printf("offset of p->left = %ld\n",
        (void *) &(p->left) - (void *)p );
}

```

An important instruction :

leq - "load effective address"
*- analogue to the "address of" operator
in C (8).*

Suppose a local variable is at -24(%rbp)
- its address is %rbp - 24
obviously!

lea -24(%rbp), %rax
- just puts %rbp - 24 into
%rax

Exactly equivalent to:

mov %rbp, %rax
sub \$24, %rax

You can use lea to take the address
of an array element : $8(a[i])$

lea (%rcx, %r8, 4), %rbx
 $\underbrace{\quad}_{\%rcx + (\%r8 \times 4)}$

lea just computes addresses, it
doesn't access memory!

mov (% rcx , % r8 , 4), % rbx

 this computes the address
and fetches the content
at that address in memory.

Suppose, in assembly, I want to call
a C function that takes a pointer
as a parameter.

void increment(int *p)

{
 (*p)++;
}

How do we call increment from assembly, passing the address of a local variable?

-g:

push %rbp

mov %rsp, %rbp

allocate space for a local variable

subq \$16, %rsp # multiple of 16

4-byte local variable will

be at -4(%rbp)

mov \$27, -4(%rbp) # x=27

want to call increment(%x)

→ `llq -4(%rbp), %rdi` # parameter
= 8x

call -increment

x would have the value 28